



Task3:-

Objective:-

- Analyze and transform large-scale multivariate time series data efficiently by:
- Implementing vectorized rolling window statistics and EWMA.
- Performing spectral analysis via FFT and filtering.
- Comparing implementations: pure NumPy, pandas built-ins, and Numba-optimized or stride-trick-accelerated variants.
- Selecting the fastest method dynamically based on input size.

Project Structure:-

```
Time_series_project/
├── timeseries_utils.py    # Core transformation functions (NumPy, pandas, Numba)
├── benchmark.py          # Benchmarks all implementations on synthetic datasets
├── benchmark_results.csv  # Benchmark timing and memory usage results
├── report.md / report.pdf # Performance comparison and summary
├── plots/                # Generated plots (speed, memory use, FFT visuals)
└── data/                 # Optionally, store large synthetic CSV datasets
```

Component 1: Rolling Window Statistics

Target Functions:-

- Rolling_mean(x, window)
- Rolling_var(x, window)

Implementation Variants:-



Variant	Notes
<code>numpy_naive</code>	Simple loop/vector approach
<code>stride_trick</code>	C-contiguous memory view with <code>as_strided</code>
<code>numba_accel</code>	Just-In-Time (JIT) accelerated with <code>@njit</code>
<code>pandas_rolling</code>	<code>df.rolling(window).mean()</code> – simple but slower at scale

✿ Component 2: Exponentially Weighted Moving Average (EWMA)

- Recursive definition:

$$y_t = \alpha x_t + (1 - \alpha)y_{t-1}$$

📌 Implementation Options:-

Variant	Notes
<code>Numpy_loop.</code>	Pure NumPy implementation (explicit recursion)



Numba_accel. JIT-accelerated for performance
Pandas_ewm df.ewm(alpha= α).mean() — clean but heavier on memory

Also include:-

•Ewma_cov(x, y, alpha): exponentially weighted covariance.

✂ Component 3: FFT Spectral Analysis & Filtering

- Use np.fft.rfft or scipy.signal for:
- Power Spectral Density (PSD): visualize energy in frequency bands.
- Band-pass filtering: retain frequencies in a defined range.

Functions:-

```
Def compute_psd(signal: np.ndarray, fs: float) -> Tuple[np.ndarray, np.ndarray]:
```

```
...
```

```
Def bandpass_filter(signal: np.ndarray, fs: float, low: float, high: float) -> np.ndarray:
```

```
...
```

🔧 Component 4: Benchmarking & Auto-Selection

- Create benchmark.py to:
- Generate synthetic time series data (1M+ rows)
- Time each implementation (e.g., time.perf_counter)
- Record memory usage (e.g., tracemalloc, psutil)
- Save results to benchmark_results.csv

Also include:-

```
Def auto_select_method(data: np.ndarray, task: str) -> str:
```



"""

Returns 'numba', 'numpy', or 'pandas' depending on performance heuristics.

"""



Visualizations (in benchmark.py)

- Runtime vs. Data Size (line plot)
- Memory usage comparison (bar plot)
- FFT spectrum plots (log scale)
- Use matplotlib or plotly.



Report Contents (report.md or .pdf)

Sections :-

- *Introduction*
 - Task overview and motivation
- *Methodology*
 - Describe algorithms for each transformation
 - Explanation of acceleration techniques (Numba, stride tricks)
- *Performance Analysis*
 - Runtime plots
 - Memory consumption tables
 - Auto-selection strategy
- *Conclusion & Recommendations*



- Best approach per task/data size
- When pandas is better vs. Custom

Dependencies:-

Bash-

Pip install numpy pandas numba matplotlib psutil

•Optional-

Bash-

Pip install plotly seaborn scipy memory_profiler

1. Starter timeseries_utils.py-

```
# timeseries_utils.py
```

```
Import numpy as np
```

```
Import pandas as pd
```

```
From numpy.lib.stride_tricks import as_strided
```

```
From numba import njit
```

```
From typing import Tuple
```

```
# -----
```

```
# Rolling Mean – NumPy naive
```

```
# -----
```

```
Def rolling_mean_numpy(x: np.ndarray, window: int) -> np.ndarray:
```

```
    Return np.convolve(x, np.ones(window)/window, mode='valid')
```

```
# -----
```

```
# Rolling Mean – Stride Trick
```

```
# -----
```

```
Def rolling_mean_stride(x: np.ndarray, window: int) -> np.ndarray:
```

```
    Shape = (x.size – window + 1, window)
```



```
Strides = (x.strides[0], x.strides[0])
Rolled = as_strided(x, shape=shape, strides=strides)
Return rolled.mean(axis=1)

# -----
# Rolling Mean – Numba JIT
# -----
@njit
Def rolling_mean_numpy(x: np.ndarray, window: int) -> np.ndarray:
    Result = np.empty(x.size - window + 1)
    For i in range(result.size):
        Result[i] = np.mean(x[i:i+window])
    Return result

# -----
# EWMA – NumPy
# -----
Def ewma_numpy(x: np.ndarray, alpha: float) -> np.ndarray:
    Result = np.zeros_like(x)
    Result[0] = x[0]
    For t in range(1, len(x)):
        Result[t] = alpha * x[t] + (1 - alpha) * result[t - 1]
    Return result

# -----
# EWMA – Numba
# -----
@njit
Def ewma_numba(x: np.ndarray, alpha: float) -> np.ndarray:
    Result = np.empty_like(x)
    Result[0] = x[0]
    For t in range(1, len(x)):
        Result[t] = alpha * x[t] + (1 - alpha) * result[t - 1]
    Return result
```

✅ 2. Ready-made benchmark.py Script-

```
# benchmark.py

Import numpy as np
Import time
Import pandas as pd
```



```
Import matplotlib.pyplot as plt
from timeseries_utils.py import (
    Rolling_mean_numpy,
    Rolling_mean_stride,
    Rolling_mean_numba,
    Ewma_numpy,
    Ewma_numba
)

Window = 50
Alpha = 0.1
Sizes = [10_000, 100_000, 1_000_000]
Methods = {
    "rolling_numpy": rolling_mean_numpy,
    "rolling_stride": rolling_mean_stride,
    "rolling_numba": rolling_mean_numba,
    "ewma_numpy": ewma_numpy,
    "ewma_numba": ewma_numba
}

Results = []

For size in sizes:
    X = np.random.rand(size)
    For name, func in methods.items():
        Try:
            T0 = time.perf_counter()
            If "rolling" in name:
                Func(x, window)
            Else:
                Func(x, alpha)
            Elapsed = time.perf_counter() - t0
            Results.append((name, size, elapsed))
        Except Exception as e:
            Results.append((name, size, np.nan))

Df = pd.DataFrame(results, columns=["Method", "DataSize", "TimeSeconds"])
Df.to_csv("benchmark_results.csv", index=False)

# Plot
Plt.figure(figsize=(10, 6))
For method in df["Method"].unique():
    Subset = df[df["Method"] == method]
    Plt.plot(subset["DataSize"], subset["TimeSeconds"], label=method)
Plt.xlabel("Data Size")
```



```
Plt.ylabel("Time (s)")
Plt.title("Benchmark: Rolling Mean & EWMA")
Plt.legend()
Plt.grid(True)
Plt.tight_layout()
Plt.savefig("plots/benchmark_plot.png")
Plt.show()
```

✅ 3. Template: report.md-

High-Performance Time Series Transformations

📄 Overview

This report evaluates several approaches to compute rolling means and exponentially weighted moving averages (EWMA) on large time series datasets using NumPy, pandas, stride tricks, and Numba.

📊 Methods Compared

Method	Backend	Optimized For
rolling_numpy	NumPy	Simplicity
rolling_stride	NumPy	Memory efficiency
rolling_numba	Numba	Speed (JIT compiled)
ewma_numpy	NumPy	Recursive updates
ewma_numba	Numba	High performance

🏎️ Benchmarking

Dataset

Synthetic univariate time series generated with:

- Sizes: 10K, 100K, 1M
- Distribution: Uniform [0, 1]

Results

![Benchmark Plot](plots/benchmark_plot.png)

📌 Observations



- **Stride tricks** outperform basic NumPy for rolling windows on medium sizes but scale poorly with massive data due to cache/memory pressure.
- **Numba** significantly speeds up recursive operations like EWMA.
- **NumPy** is sufficient for small datasets (<50K).
- **pandas** (not benchmarked yet here) is great for convenience but lags on performance.

🧠 Recommendations

- Use **Numba** for any recursive or cumulative statistics.
- Use **stride tricks** where vectorization is possible and memory alignment is good.
 - For streaming or real-time data, favor `ewma_numba()`.

📁 Files

- `timeseries_utils.py`: All methods
 - `benchmark.py`: Test runner and visualizer
 - `benchmark_results.csv`: Timings
 - `plots/benchmark_plot.png`: Performance chart
-

✅ Next Steps

- Package all this into a ZIP file?
- Add FFT analysis and band-pass filtering as the next step?



CYART

inquiry@cyart.io

www.cyart.io



CYART

inquiry@cyart.io

www.cyart.io



CYART

inquiry@cyart.io

www.cyart.io